

Cardano.BM - benchmarking and logging

Alexander Diemand

Andreas Triantafyllos

November 2018

Abstract

abstract ...

Contents

1	Cardano BM	2
1.1	Introduction	2
1.2	Overview	2
1.3	Examples	2
1.4	Code listings	2
1.4.1	Cardano.BM.Observer.STM	2
1.4.2	Cardano.BM.Observer.Monadac	4
1.4.3	BaseTrace	5
1.4.4	Cardano.BM.Trace	6
1.4.5	Cardano.BM.Controller	9
1.4.6	Cardano.BM.Counters	10
1.4.7	Cardano.BM.Counters.Dummy	11
1.4.8	Cardano.BM.Counters.Linux	11
1.4.9	Data	19
1.4.10	Cardano.BM.Agggregated	23
1.4.11	Cardano.BM.Output.Switchboard	25
1.4.12	Cardano.BM.Output.Katip	26
1.4.13	Cardano.BM.Output.Aggregation	31

Chapter 1

Cardano BM

1.1 Introduction

introduction ...

1.2 Overview

In figure 1.2 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The relationship with a triangle at one end would signify "inheritance", but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.3 Examples

examples ...

1.4 Code listings

1.4.1 Cardano.BM.Observer.STM

```
stmWithLog :: STM.STM (t,[LogObject]) → STM.STM (t,[LogObject])
stmWithLog action = action
```

```
bracketObserveIO :: Trace IO → Text → STM.STM t → IO t
bracketObserveIO logTrace0 name action = do
  (traceTransformer,logTrace) ← transformTrace name logTrace0
  bracketObserveIO' traceTransformer logTrace action
bracketObserveIO' :: TraceTransformer → Trace IO → STM.STM t → IO t
bracketObserveIO' NoTrace _ action =
  STM.atomically action
bracketObserveIO' traceTransformer logTrace action = do
  countersid ← observeOpen traceTransformer logTrace
  -- run action, returns result only
  t ← STM.atomically action
```

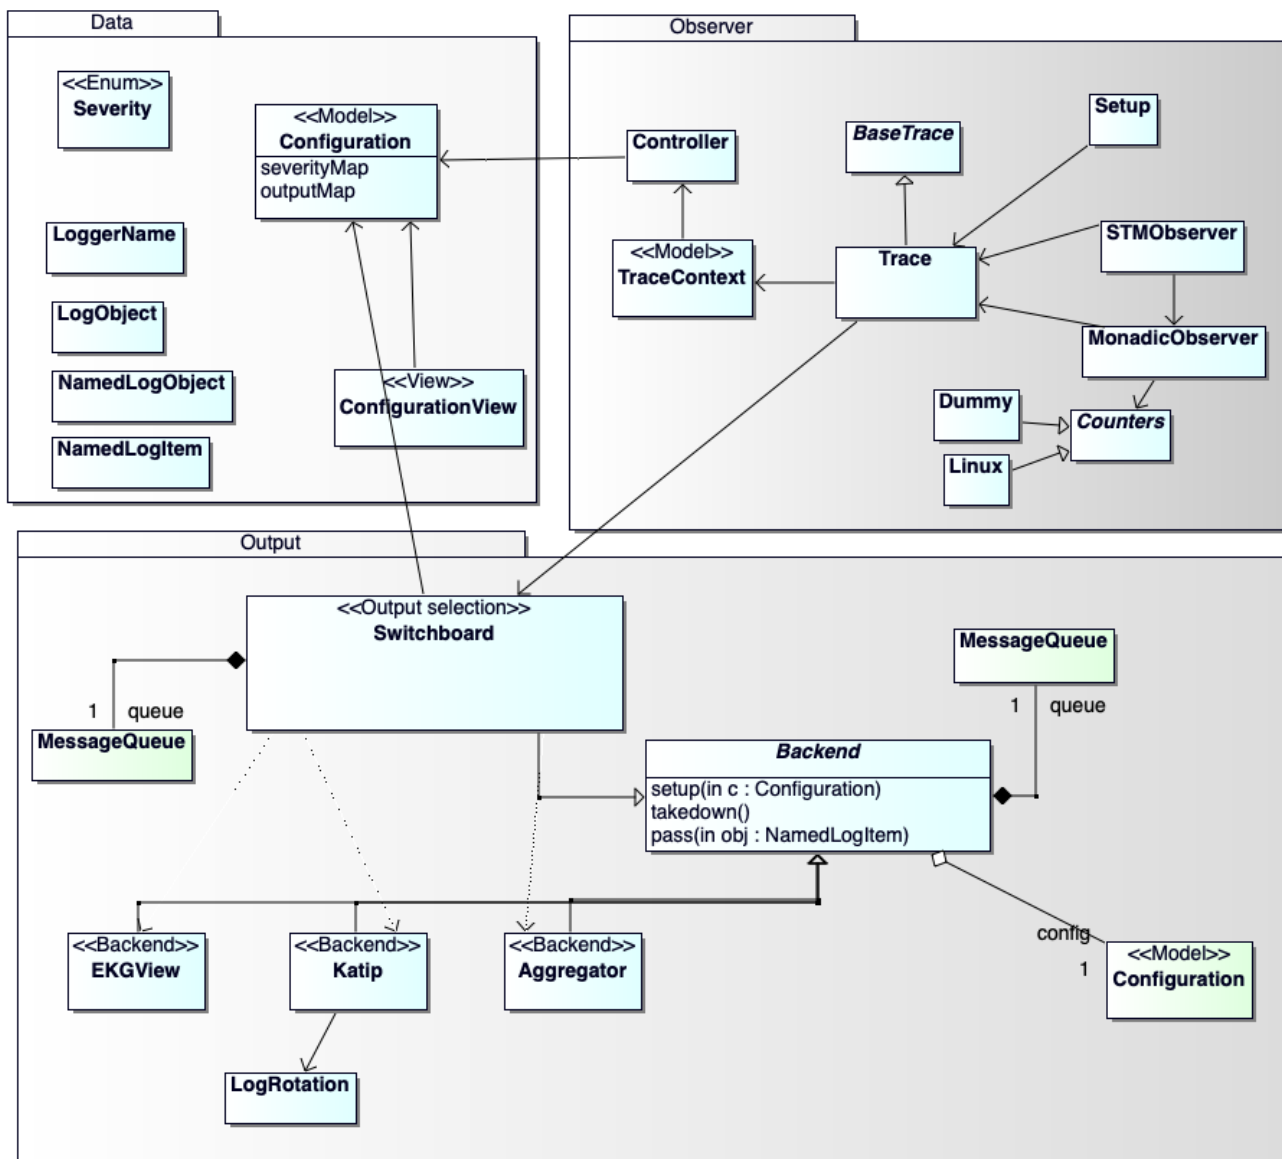


Figure 1.1: Overview of module relationships

```

    observeClose traceTransformer logTrace countersid [ ]
    pure t

bracketObserveLogIO :: Trace IO → Text → STM.STM (t,[LogObject]) → IO t
bracketObserveLogIO logTrace0 name action = do
    (traceTransformer,logTrace) ← transformTrace name logTrace0
    bracketObserveLogIO' traceTransformer logTrace action
bracketObserveLogIO' :: TraceTransformer → Trace IO → STM.STM (t,[LogObject]) → IO t
bracketObserveLogIO' NoTrace _ action = do
    (t,_) ← STM.atomically $ stmWithLog action
    pure t
bracketObserveLogIO' traceTransformer logTrace action = do
    countersid ← observeOpen traceTransformer logTrace
    -- run action, return result and log items
    (t,as) ← STM.atomically $ stmWithLog action
    observeClose traceTransformer logTrace countersid as
    pure t

```

1.4.2 Cardano.BM.Observer.Monad

```

-- Observes an action and adds name given in the logger
-- name of the given Trace. If the empty Text is
-- given as name then the logger name remains untouched.
bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = do
    (traceTransformer,logTrace) ← transformTrace name logTrace0
    bracketObserveIO' traceTransformer logTrace action
bracketObserveIO' :: TraceTransformer → Trace IO → IO t → IO t
bracketObserveIO' NoTrace _ action = action
bracketObserveIO' traceTransformer logTrace action = do
    countersid ← observeOpen traceTransformer logTrace
    -- run action
    t ← action
    observeClose traceTransformer logTrace countersid [ ]
    pure t

observeOpen :: TraceTransformer → Trace IO → IO CounterState
observeOpen traceTransformer logTrace = do
    identifier ← newUnique
    logInfo logTrace $ "Opening: " <> pack (show $ hashUnique identifier)
    -- take measurement
    counters ← readCounters traceTransformer
    let state = CounterState identifier counters
    -- send opening message to Trace

```

```

    traceNamedObject logTrace $ ObserveOpen state
    return state

```

```

observeClose :: TraceTransformer → Trace IO → CounterState → [LogObject] → IO ()
observeClose traceTransformer logTrace (CounterState identifier _) logObjects = do
    logInfo logTrace $ "Closing: " <> pack (show $ hashUnique identifier)
    -- take measurement
    counters ← readCounters traceTransformer
    -- send closing message to Trace
    traceNamedObject logTrace $ ObserveClose (CounterState identifier counters)
    -- trace the messages gathered from inside the action
    forM_ logObjects $ traceNamedObject logTrace

```

1.4.3 BaseTrace

Contravariant

A covariant is a functor: $F A \rightarrow F B$

A contravariant is a functor: $F B \rightarrow F A$

$Op\ a\ b$ implements the dual to 'arrow' " $getOp :: b \rightarrow a$ ", which when applied to a *BaseTrace* of type " $Op\ (m\ ())\ s$ ", yields " $s \rightarrow m\ ()$ ". In our case, Op accepts an action in a monad m with input type *LogNamed LogObject* (see 'Trace').

```

newtype BaseTrace m s = BaseTrace { runTrace :: Op (m ()) s }

```

contramap

A covariant functor defines the function " $fmap :: (a \rightarrow b) \rightarrow f\ a \rightarrow f\ b$ ". In case of a contravariant functor, it is the dual function " $contramap :: (a \rightarrow b) \rightarrow f\ b \rightarrow f\ a$ " which is defined.

In the following instance, $runTrace$ extracts type " $Op\ (m\ ())\ s$ " to which $contramap$ applies f , thus " $f\ s \rightarrow m\ ()$ ". The constructor *BaseTrace* restores " $Op\ (m\ ())\ (f\ s)$ ".

```

instance Contravariant (BaseTrace m) where
    contramap f = BaseTrace ∘ contramap f ∘ runTrace

```

traceWith

Accepts a *Trace* and some payload s . First it gets the contravariant from the *Trace* as type " $Op\ (m\ ())\ s$ " and, after " $getOp :: b \rightarrow a$ " which translates to " $s \rightarrow m\ ()$ ", calls the action on the *LogNamed LogObject*.

```

traceWith :: BaseTrace m s → s → m ()
traceWith = getOp ∘ runTrace

```

natTrace

Natural transformation from monad m to monad n .

```
natTrace :: (forall x o m x → n x) → BaseTrace m s → BaseTrace n s
natTrace nat (BaseTrace (Op tr)) = BaseTrace $ Op $ nat o tr
```

noTrace

A *Trace* that discards all inputs.

```
noTrace :: Applicative m ⇒ BaseTrace m a
noTrace = BaseTrace $ Op $ const (pure ())
```

1.4.4 Cardano.BM.Trace

```
appendName :: MonadIO m ⇒ LoggerName → Trace m → m (Trace m)
appendName name (ctx, trace) = do
  -- append the name to the existing one
  let previousLoggerName = loggerName ctx
      newLoggerName = appendWithDot previousLoggerName name
      ctx' = ctx {loggerName = newLoggerName}
  -- inherit the Severity of the given Trace.
  mPreviousSeverity ← liftIO $ getNamedSeverity ctx' previousLoggerName
  case mPreviousSeverity of
    Nothing → return (ctx', trace)
    Just sev → liftIO $ setNamedSeverity ctx newLoggerName sev
                >> return (ctx', trace)

  -- return a BaseTrace from a TraceNamed
named :: BaseTrace m (LogNamed i) → LoggerName → BaseTrace m i
named trace name = contramap (LogNamed name) trace
```

TODO remove *locallock*

```
locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()
```

Concrete Trace on stdout

This function returns a trace with an action of type $"(LogNamed LogObject) \rightarrow IO ()"$ which will output a text message as text and all others as JSON encoded representation to the console.

```
stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace $ Op $ \lognamed →
  case lnItem lognamed of
    LP (LogMessage logItem) →
      withMVar locallock $ \_ →
```



```

        output (lnName lognamed) $ liPayload logItem
    obj →
        withMVar locallock $ \_ →
            output (lnName lognamed) $ toStrict (encodeToLazyText obj)
where
    output nm msg = TIO.putStrLn $ nm <> " :: " <> msg

```

Trace into katip's queue

```

katipTrace :: {-MVar LoggingHandler -} TraceNamed IO
katipTrace = BaseTrace $ Op $ \lognamed → do
    lh ← readMVar Internal.loggingHandler
    mayEnv ← Internal.getLogEnv lh
    case mayEnv of
        Nothing → error "logging not yet initialized. Abort."
        Just env → logItem'
            lognamed
            (K.Namespace (T.split (≡ ' . ') (lnName lognamed)))
            env
        Nothing
            (Internal.sev2klog Info)
            (K.logStr (" " :: Text))

```

Every *Trace* ends in the switchboard which then takes care of dispatching the messages to outputs

```

mainTrace :: TraceNamed IO
mainTrace = BaseTrace $ Op $ \lognamed → do
    Switchboard.pass lognamed

```

Concrete Trace into a TVar

```

traceInTVar :: STM.TVar [a] → BaseTrace STM.STM a
traceInTVar tvar = BaseTrace $ Op $ \a → STM.modifyTVar tvar ((:) a)
traceInTVarIO :: STM.TVar [LogObject] → TraceNamed IO
traceInTVarIO tvar = BaseTrace $ Op $ \ln →
    STM.atomically $ STM.modifyTVar tvar ((:) (lnItem ln))
traceNamedInTVarIO :: STM.TVar [LogNamed LogObject] → TraceNamed IO
traceNamedInTVarIO tvar = BaseTrace $ Op $ \ln →
    STM.atomically $ STM.modifyTVar tvar ((:) ln)

traceConditionally
    :: (MonadIO m)
    ⇒ TraceContext → BaseTrace m LogObject → LogObject
    → m ()

```

```

traceConditionally ctx logTrace msg@(LP (LogMessage item)) = do
  flag ← liftIO $ checkSeverity ctx item
  when flag $ traceWith logTrace msg
traceConditionally _ logTrace logObject = traceWith logTrace logObject

```

Enter message into a trace

The function `traceNamedItem` creates a `LogObject` and threads this through the action defined in the `Trace`.

```

traceNamedItem
  :: (MonadIO m)
  ⇒ Trace m
  → LogSelection
  → Severity
  → Text
  → m ()
traceNamedItem (ctx, logTrace) p s m =
  let logmsg = LP $ LogMessage $ LogItem { liSelection = p
    , liSeverity = s
    , liPayload = m
  }
  in
    traceConditionally ctx (named logTrace (loggerName ctx)) $ logmsg
logDebug, logInfo, logNotice, logWarning, logError
  :: (MonadIO m) ⇒ Trace m → Text → m ()
logDebug logTrace = traceNamedItem logTrace Both Debug
logInfo logTrace   = traceNamedItem logTrace Both Info
logNotice logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError logTrace  = traceNamedItem logTrace Both Error
logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS
  :: (MonadIO m) ⇒ Trace m → Text → m ()
logDebugS logTrace = traceNamedItem logTrace Private Debug
logInfoS logTrace   = traceNamedItem logTrace Private Info
logNoticeS logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS logTrace  = traceNamedItem logTrace Private Error
logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP
  :: (MonadIO m) ⇒ Trace m → Text → m ()
logDebugP logTrace = traceNamedItem logTrace Public Debug
logInfoP logTrace   = traceNamedItem logTrace Public Info
logNoticeP logTrace = traceNamedItem logTrace Public Notice
logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP logTrace  = traceNamedItem logTrace Public Error
logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP
  :: (MonadIO m) ⇒ Trace m → Text → m ()

```

```

logDebugUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP logTrace  = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP logTrace  = traceNamedItem logTrace PublicUnsafe Error

```

```

traceNamedObject
  :: Trace m
  → LogObject
  → m ()
traceNamedObject (ctx, logTrace) = traceWith (named logTrace (loggerName ctx))

```

transformTrace

```

-- Transforms the Trace given according to content of
-- TraceTransformerMap using the logger name of the
-- current Trace appended with the given name. If the
-- empty Text is given as name then the logger name
-- remains untouched.
transformTrace :: MonadIO m ⇒ Text → Trace m → m (TraceTransformer, Trace m)
transformTrace name tr@(ctx, _) = do
  traceTransformer ← liftIO $ findTraceTransformer tr $ appendWithDot (loggerName ctx) name
  case traceTransformer of
    Neutral      → do
      tr' ← appendName name tr
      return $ (traceTransformer, tr')
    UntimedTrace → do
      tr' ← appendName name tr
      return $ (traceTransformer, tr')
    NoTrace      → return (traceTransformer, (ctx, BaseTrace $ Op $ \_ → pure ()))
    DropOpening  → return (traceTransformer, (ctx, BaseTrace $ Op $ \lognamed →
      case lnItem lognamed of
        ObserveOpen _ → return ()
        obj            → traceNamedObject tr obj))
    ObservableTrace _ → do
      tr' ← appendName name tr
      return $ (traceTransformer, tr')

```

1.4.5 Cardano.BM.Controller

```

findTraceTransformer :: Trace m → Text → IO TraceTransformer
findTraceTransformer (ctx, _) name =
  withMVar (controller ctx) $ λtc →
    return $ findWithDefault Neutral name (traceTransformers tc)
appendWithDot :: LoggerName → LoggerName → LoggerName

```

```

appendWithDot "" newName = take 50 newName
appendWithDot xs "" = xs
appendWithDot xs newName = take 50 $ xs <> "." <> newName
insertInController :: Monad m => Trace m → Text → TraceTransformer → IO ()
insertInController (ctx, _) name trans = do
  let currentLoggerName = loggerName ctx
      name' = appendWithDot currentLoggerName name
  modifyMVar_ (controller ctx) $ λtc →
    return $ tc { traceTransformers = insert name' trans (traceTransformers tc) }
setMinSeverity :: Trace m → Severity → IO ()
setMinSeverity (ctx, _) newMinSeverity =
  modifyMVar_ (controller ctx) $ λtc →
    return $ tc { minSeverity = newMinSeverity }
setNamedSeverity :: TraceContext → LoggerName → Severity → IO ()
setNamedSeverity ctx name newSeverity =
  modifyMVar_ (controller ctx) $ λtc →
    return $ tc { severityMap = insert name newSeverity (severityMap tc) }
getNamedSeverity :: TraceContext → LoggerName → IO (Maybe Severity)
getNamedSeverity ctx name = withMVar (controller ctx) $ λtc →
  return $ lookup name (severityMap tc)

```

checkSeverity

```

checkSeverity :: TraceContext → LogItem → IO Bool
checkSeverity ctx item = do
  let name = loggerName ctx
      itemSev = liSeverity item
  withMVar (controller ctx) $ λtc → do
    let globalSev = minSeverity tc
    case lookup name (severityMap tc) of
      Nothing → return (itemSev ≥ globalSev)
      Just specificSev → return ((itemSev ≥ globalSev) ∧ (itemSev ≥ specificSev))

```

1.4.6 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.

Currently, we only support *Linux* with its 'proc' filesystem.

```

{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
(
  Platform.readCounters

```

```

    ) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif

```

1.4.7 Cardano.BM.Counters.Dummy

This is a dummy definition of `readCounters` on platforms that do not support the 'proc' filesystem from which we would read the counters.

we could well imagine that some day we support all platforms

```

module Cardano.BM.Counters.Dummy
(
    readCounters
) where
import Cardano.BM.Data (Counter, TraceTransformer (..))
readCounters :: TraceTransformer → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace _) = return []

```

1.4.8 Cardano.BM.Counters.Linux

```

nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ ('div' 1000)

```

we have to expand the `getMonoClock` and `readMemStats` functions with ones that read full data from `proc`

```

readCounters :: TraceTransformer → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace tts) = foldrM (λ(sel,fun) a →
    if sel 'member' tts
    then (fun >>= λxs → return $ a ++ xs)
    else return a) [] selectors
where
    selectors = [(MonotonicClock, getMonoClock)
                , (MemoryStats, readProcStatM)

```

```

    ,(ProcessStats,readProcStats)
    ,(IOStats,readProcIO)
  ]
getMonoClock :: IO [Counter]
getMonoClock = do
  t ← getMonotonicTimeNSec
  return [MonotonicClockTime "monoclock" $ nominalTimeToMicroseconds t]

```

```

pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc </> (show pid) </> "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc </> (show pid) </> "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc </> (show pid) </> "io"

```

Reading from a file in /proc/<pid>

```

readProcList :: FilePath → IO [Integer]
readProcList fp = do
  cs ← readFile fp
  return $ map (\s → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)

```

readProcStatM - /proc/<pid>/statm

/proc/[pid]/statm

Provides information about memory usage, measured in pages. The columns are:

size	(1) total program size (same as VmSize in /proc/[pid]/status)
resident	(2) resident set size (same as VmRSS in /proc/[pid]/status)
shared	(3) number of resident shared pages (i.e., backed by a file) (same as RssFile+RssShmem in /proc/[pid]/status)
text	(4) text (code)
lib	(5) library (unused since Linux 2.6; always 0)
data	(6) data + stack
dt	(7) dirty pages (unused since Linux 2.6; always 0)

```

readProcStatM :: IO [Counter]
readProcStatM = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcStatM pid)
  ps ← return $ zip colnames ps0
  forM ps (\(n,i) → return $ MemoryCounter n i)

```

where

```
colnames::[Text]
colnames = ["size", "resident", "shared", "text", "unused", "data", "unused"]
```

readProcStats - //proc//<pid >//stat

```
/proc/[pid]/stat
```

Status information about the process. This is used by ps(1). It is defined in fs/proc/array.c.

The fields, in order, with their proper scanf(3) format specifiers, are listed below. Whether or not certain of these fields display valid information is governed by the PT_TRACE_MODE_READ_FSCREDS | PT_TRACE_MODE_NOAUDIT check (refer to ptrace(2)). If the check fails, then the field value is displayed as 0. The affected fields are indicated with an asterisk.

(1) pid %d

The process ID.

(2) comm %s

The filename of the executable, in parentheses. This is visible only if the process's executable is swapped out.

(3) state %c

One of the following characters, indicating process state:

R Running

S Sleeping in an interruptible wait

D Waiting in uninterruptible disk sleep

Z Zombie

T Stopped (on a signal) or (before Linux 2.6.33) trace stopped

t Tracing stop (Linux 2.6.33 onward)

W Paging (only before Linux 2.6.0)

X Dead (from Linux 2.6.0 onward)

x Dead (Linux 2.6.33 to 3.13 only)

K Wakekill (Linux 2.6.33 to 3.13 only)

W Waking (Linux 2.6.33 to 3.13 only)

P Parsed (Linux 3.9 to 3.13 only)

(4) ppid %d

The PID of the parent of this process.

(5) pgrp %d

The process group ID of the process.

(6) session %d

The session ID of the process.

(7) tty_nr %d

The controlling terminal of the process. (The minor device number is in bits 31 to 20 and 7 to 0; the major device number is in bits 31 to 20).

(8) tpgid %d

The ID of the foreground process group of the controlling terminal of the process.

(9) flags %u

The kernel flags word of the process. For bit meanings, see the kernel source file include/linux/sched.h. Details depend on the kernel version.

The format for this field was %lu before Linux 2.6.

(10) minflt %lu

The number of minor faults the process has made which have not required a swap from disk.

(11) cminflt %lu

The number of minor faults that the process's waited-for children have made.

(12) majflt %lu

The number of major faults the process has made which have required a swap from disk.

(13) cmajflt %lu

The number of major faults that the process's waited-for children have made.

(14) utime %lu

Amount of time that this process has been scheduled in user mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). This includes guest time, `guest_time` (when running on a virtual CPU, see below), so that applications that are not aware of virtualization do not lose that time from their calculations.

(15) stime %lu

Amount of time that this process has been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

- (16) `cutime %ld`
 Amount of time that this process's waited-for children have been scheduled in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). (See also `cguest_time` (time spent running a virtual CPU, see below).
- (17) `cstime %ld`
 Amount of time that this process's waited-for children have been scheduled in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (18) `priority %ld`
 (Explanation for Linux 2.6) For processes running a real-time scheduling policy (see `sched_setscheduler(2)`), this is the negated scheduling priority; is, a number in the range -2 to -100, corresponding to real-time priorities. For processes running under a non-real-time scheduling policy, this is the nice value as represented in the kernel. The kernel stores nice values in the range 0 (high) to 39 (low), corresponding to the user-visible nice values.
- (19) `nice %ld`
 The nice value (see `setpriority(2)`), a value in the range 19 (low priority) to 0 (high priority).
- (20) `num_threads %ld`
 Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was coded to 0 as a placeholder for an earlier removed field.
- (21) `itrealvalue %ld`
 The time in jiffies before the next SIGALRM is sent to the process or it terminates. Since kernel 2.6.17, this field is no longer maintained, and is hard-coded to 0.
- (22) `starttime %llu`
 The time the process started after system boot. In kernels before Linux 2.6, this value is expressed in jiffies. Since Linux 2.6, the value is expressed in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- The format for this field was `%lu` before Linux 2.6.
- (23) `vsize %lu`
 Virtual memory size in bytes.
- (24) `rss %ld`
 Resident Set Size: number of pages the process has in real memory. This includes pages which count toward text, data, or stack space. This does not include pages which have been demand-loaded in, or which are swapped out.
- (25) `rsslim %lu`
 Current soft limit in bytes on the rss of the process; see the description of `getrlimit(2)`.

- (26) startcode %lu [PT]
The address above which program text can run.
- (27) endcode %lu [PT]
The address below which program text can run.
- (28) startstack %lu [PT]
The address of the start (i.e., bottom) of the stack.
- (29) kstkesp %lu [PT]
The current value of ESP (stack pointer), as found in the kernel stack.
- (30) kstkeip %lu [PT]
The current EIP (instruction pointer).
- (31) signal %lu
The bitmap of pending signals, displayed as a decimal number. Obsolete; use /proc/[pid]/status instead to provide information on real-time signals; use /proc/[pid]/status instead for real-time signals.
- (32) blocked %lu
The bitmap of blocked signals, displayed as a decimal number. Obsolete; use /proc/[pid]/status instead to provide information on real-time signals; use /proc/[pid]/status instead for real-time signals.
- (33) sigignore %lu
The bitmap of ignored signals, displayed as a decimal number. Obsolete; use /proc/[pid]/status instead to provide information on real-time signals; use /proc/[pid]/status instead for real-time signals.
- (34) sigcatch %lu
The bitmap of caught signals, displayed as a decimal number. Obsolete; use /proc/[pid]/status instead to provide information on real-time signals; use /proc/[pid]/status instead for real-time signals.
- (35) wchan %lu [PT]
This is the "channel" in which the process is waiting. It is the kernel where the process is sleeping. The corresponding symbolic name is /proc/[pid]/wchan.
- (36) nswap %lu
Number of pages swapped (not maintained).
- (37) cnswap %lu
Cumulative nswap for child processes (not maintained).
- (38) exit_signal %d (since Linux 2.1.22)
Signal to be sent to parent when we die.
- (39) processor %d (since Linux 2.2.8)
CPU number last executed on.

- (40) `rt_priority %u` (since Linux 2.5.19)
Real-time scheduling priority, a number in the range 1 to 99 for processes with real-time policy, or 0, for non-real-time processes (see `sched_setscheduler(2)`).
- (41) `policy %u` (since Linux 2.5.19)
Scheduling policy (see `sched_setscheduler(2)`). Decode using the `linux/sched.h` header.
- The format for this field was `%lu` before Linux 2.6.22.
- (42) `delayacct_blkio_ticks %llu` (since Linux 2.6.18)
Aggregated block I/O delays, measured in clock ticks (centiseconds).
- (43) `guest_time %lu` (since Linux 2.6.24)
Guest time of the process (time spent running a virtual CPU for a guest), measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (44) `cguest_time %ld` (since Linux 2.6.24)
Guest time of the process's children, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (45) `start_data %lu` (since Linux 3.3) [PT]
Address above which program initialized and uninitialized (BSS) data starts.
- (46) `end_data %lu` (since Linux 3.3) [PT]
Address below which program initialized and uninitialized (BSS) data ends.
- (47) `start_brk %lu` (since Linux 3.3) [PT]
Address above which program heap can be expanded with `brk(2)`.
- (48) `arg_start %lu` (since Linux 3.5) [PT]
Address above which program command-line arguments (`argv`) are placed.
- (49) `arg_end %lu` (since Linux 3.5) [PT]
Address below program command-line arguments (`argv`) are placed.
- (50) `env_start %lu` (since Linux 3.5) [PT]
Address above which program environment is placed.
- (51) `env_end %lu` (since Linux 3.5) [PT]
Address below which program environment is placed.
- (52) `exit_code %d` (since Linux 3.5) [PT]
The thread's exit status in the form reported by `waitpid(2)`.

```
readProcStats :: IO [Counter]
readProcStats = do
```

```

pid ← getProcessID
ps0 ← readProcList (pathProcStat pid)
ps ← return $ zip colnames ps0
forM ps (λ(n,i) → return $ StatInfo n i)
where
  colnames :: [Text]
  colnames = [ "pid", "unused", "unused", "ppid", "pgrp", "session", "ttynr", "tpgid", "flags", "minflt",
    , "cminflt", "majflt", "cmajflt", "utime", "stime", "cutime", "cstime", "priority", "nice", "num",
    , "itrealvalue", "starttime", "vsize", "rss", "rsslim", "startcode", "endcode", "startstack", "proc",
    , "signal", "blocked", "sigignore", "sigcatch", "wchan", "nswap", "cnsnap", "exitsignal", "proc",
    , "policy", "blkio", "guesttime", "cguesttime", "startdata", "enddata", "startbrk", "argstart",
    , "envend", "exitcode"
  ]

```

readProcIO - //proc//<pid >//io

/proc/[pid]/io (since kernel 2.6.20)

This file contains I/O statistics for the process, for example:

```

# cat /proc/3828/io
rchar: 323934931
wchar: 323929600
syscr: 632687
syscw: 632675
read\_bytes: 0
write\_bytes: 323932160
cancelled\_write\_bytes: 0

```

The fields are as follows:

rchar: characters read

The number of bytes which this task has caused to be read from storage. of bytes which this process passed to read(2) and similar system calls. as terminal I/O and is unaffected by whether or not actual physical disk read might have been satisfied from pagecache).

wchar: characters written

The number of bytes which this task has caused, or shall cause to be v caveats apply here as with rchar.

syscr: read syscalls

Attempt to count the number of read I/O operations-that is, system calls pread(2).

syscw: write syscalls

Attempt to count the number of write I/O operations-that is, system ca pwrite(2).

```
\usepackage{verbatim}
  read\_bytes: bytes read
    Attempt to count the number of bytes which this process really did cause
    storage layer. This is accurate for block-backed filesystems.

  write\_bytes: bytes written
    Attempt to count the number of bytes which this process caused to be ser

  cancelled\_write\_bytes:
    The big inaccuracy here is truncate. If a process writes 1MB to a f
    file, it will in fact perform no writeout. But it will have been account
    of write. In other words: this field represents the number of bytes w
    to not happen, by truncating pagecache. A task can cause "negative" I,
    truncates some dirty pagecache, some I/O which another task has be
    write\_bytes) will not be happening.
```

Note: In the current implementation, things are a bit racy on 32-bit systems
process B's /proc/[pid]/io while process B is updating one of these 64-bit c
see an intermediate result.

Permission to access this file is governed by a ptrace access mode PTRACE_MODE
ptrace(2).

```
readProcIO :: IO [Counter]
readProcIO = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcIO pid)
  ps ← return $ zip colnames ps0
  forM ps (\(n,i) → return $ IOCounter n i)
  where
    colnames :: [Text]
    colnames = [ "rchar", "wchar", "syscr", "syscw", "rbytes", "wbytes", "cxwbytes" ]
```

1.4.9 Data

type aliases and empty types

```
type NamedLogItem = LogNamed LogObject
```

Trace

A *Trace* consists of a *TraceContext* and a *TraceNamed* in *m*.

```
type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A *TraceNamed* is a trace of type *LogNamed* with payload *LogObject*.

```
type TraceNamed m = BaseTrace m (LogNamed LogObject)
```

LogObject

```
data LogPrims = LogMessage LogItem
  | LogValue Text Integer
  deriving (Generic, Show, ToJSON)
data LogObject = LP LogPrims
  | ObserveOpen CounterState
  | ObserveClose CounterState
  deriving (Generic, Show, ToJSON)
```

TraceTransformer

```
data TraceTransformer = Neutral
  | UntimedTrace
  | NoTrace
  | DropOpening
  | ObservableTrace (Set ObservableInstance)
  deriving (Show)
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
  deriving (Eq, Ord, Show)
```

TODO *lnName* :: Text
storing a concatenation of names might be cheaper than rebuilding it for every log message

LogNamed

A *LogNamed* contains of a list of context names and some log item.

```
-- Attach a 'ContextName' to something.
data LogNamed item = LogNamed
  { lnName :: LoggerName
  , lnItem :: item
  } deriving (Show)
deriving instance Generic item => Generic (LogNamed item)
deriving instance (ToJSON item, Generic item) => ToJSON (LogNamed item)
-- Attach a 'ContextName' and Katip related info to something.
```

```

data LogNamedPlus item = LogNamedPlus
  { lnpName :: LoggerName
  , lnpItem :: item
  } deriving (Show)

```

LogItem

```

TODO liPayload :: ToObject

```

```

-- log item
data LogItem = LogItem
  { liSelection :: LogSelection
  , liSeverity :: Severity
  , liPayload :: Text -- TODO should become ToObject
  } deriving (Show, Generic, ToJSON)

-- output selection
data LogSelection =
  Public -- only to public logs.
  | PublicUnsafe -- only to public logs, not console.
  | Private -- only to private logs.
  | Both -- to public and private logs.
  deriving (Show, Generic, ToJSON)

-- severity of log message
data Severity = Debug | Info | Warning | Notice | Error
  deriving (Show, Eq, Ord, Generic, ToJSON)

instance FromJSON Severity where
  parseJSON = withText "severity" $ \case
    "Debug"    → pure Debug
    "Info"     → pure Info
    "Notice"   → pure Notice
    "Warning"  → pure Warning
    "Error"    → pure Error
    _          → pure Info -- catch all

```

Observable

```

data Counter = MonotonicClockTime Text Microsecond
  | MemoryCounter Text Integer
  | StatInfo Text Integer
  | IOCounter Text Integer
  | CpuCounter Text Integer
  deriving (Show, Generic, ToJSON)

instance ToJSON Microsecond where
  toJSON = toJSON ∘ toMicroseconds
  toEncoding = toEncoding ∘ toMicroseconds

```

```

data CounterState = CounterState {
  csIdentifier :: Unique
  ,csCounters :: [Counter]
}
deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ◦ hashUnique
  toEncoding = toEncoding ◦ hashUnique
instance Show CounterState where
  show cs = (show ◦ hashUnique) (csIdentifier cs)
    <> " => " <> (show $ csCounters cs)

```

TraceContext

```

type LoggerName = Text
data TraceContext = TraceContext {
  loggerName :: LoggerName
  ,controller :: MVar TraceController
}
type TraceTransformerMap = Map LoggerName TraceTransformer
type SeverityMap = Map LoggerName Severity
data TraceController = TraceController {
  traceTransformers :: TraceTransformerMap
  ,severityMap :: SeverityMap
  ,minSeverity :: Severity
}

```

TraceConfiguration

```

data TraceConfiguration = TraceConfiguration
  { tcOutputKind :: OutputKind
  , tcName :: LoggerName
  , tcTraceTransformer :: TraceTransformer
  , tcSeverity :: Severity
  }
data OutputKind = StdOut
  | TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  | Null
deriving Eq
diffTimeObserved :: CounterState → CounterState → Microsecond
diffTimeObserved (CounterState _ startCounters) (CounterState _ endCounters) =
  let
    startTime = getMonotonicTime startCounters

```



```

    endTime = getMonotonicTime endCounters
  in
    endTime - startTime
  where
    getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
      [(MonotonicClockTime _ micros)] → micros
      _ → error "Exactly one time measurements was expected!"
    isMonotonicClockCounter :: Counter → Bool
    isMonotonicClockCounter (MonotonicClockTime _) = True
    isMonotonicClockCounter _ = False

```

Backend

A backend is referenced through the function *pass'* which accepts a ??.

```

data Backend = Backend {pass' :: NamedLogItem → IO ()}

```

ScribeKind

This identifies katip's scribes by type.

```

data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  | DevNullSK
  deriving (Eq, Show)

```

BackendKind

This identifies the backends that can be attached to the ??.

```

data BackendKind = AggregationBK
  | EKGViewBK
  | KatipBK
  | DevNullBK
  deriving (Eq, Show)

```

1.4.10 Cardano.BM.Agregated

```

module Cardano.BM.Agregated
(
  Aggregated (..)
  , Stats (..)
  , updateAggregation
) where

```

```

data Stats = Stats {
  fmin :: Integer,
  fmax :: Integer,
  fcount :: Integer,
  fsum_A :: Integer,
  fsum_B :: Integer
} deriving (Show, Eq)

data Aggregated = Aggregated {
  fstats :: Stats,
  flast :: Integer,
  fdelta :: Stats
} deriving (Show, Eq)

```

Update aggregation

We distinguish an uninitialized from an already initialized aggregation:

```

updateAggregation :: Integer → Maybe Aggregated → Maybe Aggregated
updateAggregation v Nothing =
  Just $
    Aggregated {fstats = Stats {
      fmin = v, fmax = v, fcount = 1
      , fsum_A = v, fsum_B = v * v
    }, flast = v
    , fdelta = Stats {
      fmin = 0, fmax = 0, fcount = 0
      , fsum_A = 0, fsum_B = 0
    }
    }
updateAggregation v (Just (Aggregated (Stats _min _max _count _sumA _sumB)
  _last
  (Stats _dmin _dmax _dcount _dsumA _dsumB)
  )) =
  let delta = v - _last
  in
  Just $
    Aggregated {fstats = Stats {
      fmin = (min _min v)
      , fmax = (max _max v)
      , fcount = (_count + 1)
      , fsum_A = (_sumA + v)
      , fsum_B = (_sumB + v * v)
    }
    , flast = v
    , fdelta = Stats {
      fmin = (min _dmin delta)
      , fmax = (max _dmax delta)
      , fcount = (_dcount + 1)
    }
    }

```

```

    ,fsum_A = (.dsumA + delta)
    ,fsum_B = (.dsumB + delta * delta)
  }
}

```

1.4.11 Cardano.BM.Output.Switchboard

State representation

The switchboard is a singleton.

```

-- internal access to the switchboard
{-# NOINLINE switchboard #-}
switchboard :: MVar SwitchboardInternal
switchboard = unsafePerformIO $ do
  newMVar $ error "Switchboard MVar is not initialized."
-- Our internal state
data SwitchboardInternal = SwitchboardInternal
  { sbQueue :: TBQ.TBQueue (Maybe NamedLogItem)
  , sbDispatch :: Async.Async ()
  , sbBackends :: [Backend]
  }

```

Starting the switchboard from configuration

The queue is initialized and the message dispatcher launched. TODO: the backends should be connected according to configuration.

```

setup :: Configuration → IO ()
setup _ = do
  _ ← takeMVar switchboard
  q ← atomically $ TBQ.newTBQueue 2048
  d ← spawnDispatcher q
  -- TODO connect backends according to configuration
  let be = [Backend {pass' = Katip.pass "StdoutSK"}]
  putMVar switchboard $ SwitchboardInternal q d be
spawnDispatcher :: TBQ.TBQueue (Maybe NamedLogItem) → IO (Async.Async ())
spawnDispatcher queue = Async.async qProc
where
  qProc = do
    nli' ← atomically $ TBQ.readTBQueue queue
    case nli' of
      Just nli → do
        putStrLn $ "dispatcher read: " ++ (show nli)
        withMVar switchboard $ λsb →
          forM_ (sbBackends sb) (dispatch nli)
      _ → qProc

```

```

    Nothing → return () -- end dispatcher
    dispatch :: NamedLogItem → Backend → IO ()
    dispatch nli backend = (pass' backend) nli

```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher.

```

pass :: NamedLogItem → IO ()
pass item = do
    putStrLn $ "Cardano.BM.Output.Switchboard.pass " ++ (show item)
    withMVar switchboard $ \sb →
        atomically $ writequeue (sbQueue sb) item
where
    writequeue :: TBQ.TBQueue (Maybe NamedLogItem) → NamedLogItem → STM ()
    writequeue q i = do
        nocapacity ← TBQ.isFullTBQueue q
        if ¬ nocapacity
        then TBQ.writeTBQueue q (Just i)
        else return ()

```

Halting the switchboard

The queue is flushed before the dispatcher terminates.

```

takedown :: IO ()
takedown = do
    (q,d) ← withMVar switchboard $ \sb →
        return (sbQueue sb, sbDispatch sb)
    -- send terminating item to the queue
    atomically $ TBQ.writeTBQueue q Nothing
    -- wait for the dispatcher to exit
    _ ← Async.waitCatch d
    return ()

```

1.4.12 Cardano.BM.Output.Katip

Katip is a singleton.

```

-- internal access to katip
{-# NOINLINE katip #-}
katip :: MVar KatipInternal
katip = unsafePerformIO $ do
    newMVar $ error "Katip MVar is not initialized."
-- Our internal state
data KatipInternal = KatipInternal
    { kLogEnv :: K.LogEnv }

```

Setup *katip* and its scribes according to the configuration

```

setup :: Config.Configuration → IO ()
setup config = do
  setDefaultBackends [ Backend {pass' = Cardano.BM.Output ∘ Katip.pass (pack (show StdoutSK))}
    , Backend {pass' = Cardano.BM.Output ∘ Katip.pass (pack (show FileTextSK))}
    , Backend {pass' = Cardano.BM.Output ∘ Katip.pass (pack (show FileJsonSK))}
  ]
  cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
  -- TODO setup katip
  le0 ← K.initLogEnv
    (K.Namespace [ "ouroboros-bm" ])
    (fromString $ (unpack cfoKey) <> ":" <> showVersion mockVersion)
  -- request a new time 'getCurrentTime' at most 100 times a second
  timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
  let le1 = updateEnv le0 timer
  stdoutScribe ← mkStdoutScribeJson K.V0
  le ← register [(StdoutSK, "stdout", stdoutScribe)] le1
  _ ← takeMVar katip
  putMVar katip $ KatipInternal le
where
  updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
  updateEnv le timer =
    le {K._logEnvTimer = timer, K._logEnvHost = "hostname"}
  register :: [(ScribeKind, Text, K.Scribe)] → K.LogEnv → IO K.LogEnv
  register [] le = return le
  register ((kind, name, scribe) : scs) le =
    let name' = pack (show kind) <> ":" <> name in
    register scs ≡ K.registerScribe name' scribe scribeSettings le
  mockVersion :: Version
  mockVersion = Version [0, 1, 0, 0] []
  scribeSettings :: KC.ScribeSettings
  scribeSettings = KC.ScribeSettings bufferSize
  where
    bufferSize = 5000 -- size of the queue (in log items)

example :: IO ()
example = do
  config ← Config.setup "from_some_path.yaml"
  setup config
  pass (pack (show StdoutSK)) $ LogNamed
    {lnName = "test"
    ,lnItem = LP $ LogMessage $ LogItem
      {liSelection = Both
      ,liSeverity = Info
      ,liPayload = "Hello!"
      }
    }

```

```

pass (pack (show StdoutSK)) $ LogNamed
  {lnName = "test"
  ,lnItem = LP $ LogValue "cpu-no" 1
  }

-- useful instances for Katip
deriving instance K.ToObject LogObject
deriving instance K.ToObject LogItem
deriving instance K.ToObject (Maybe LogObject)
instance KC.LogItem LogObject where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem LogItem where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem (Maybe LogObject) where
  payloadKeys _ = KC.AllKeys

pass :: Text → NamedLogItem → IO ()
pass backend namedLogItem = withMVar katip $ λk → do
  -- TODO go through list of registered scribes
  -- and put into queue of scribe if backend kind matches
  -- compare start of name of scribe to (show backend <> "::")
  let env = kLogEnv k
  forM_ (Map.toList $ K._logEnvScribes env) $
    λ(scName, (KC.ScribeHandle _ shChan)) →
      -- check start of name to match ScribeKind
      if backend `isPrefixOf` scName
      then do
        let item = lnItem namedLogItem
        let (sev, msg, payload) = case item of
          (LP (LogMessage logItem)) →
            (liSeverity logItem, liPayload logItem, Nothing)
          _ → (Info, "", Just item)
        threadIdText ← KC.mkThreadIdText <$> myThreadId
        let ns = lnName namedLogItem
        itemTime ← env ^. KC.logEnvTimer
        let itemKatip = K.Item {
          _itemApp      = env ^. KC.logEnvApp
          , _itemEnv     = env ^. KC.logEnvEnv
          , _itemSeverity = sev2klog sev
          , _itemThread  = threadIdText
          , _itemHost     = env ^. KC.logEnvHost
          , _itemProcess  = env ^. KC.logEnvPid
          , _itemPayload  = payload
          , _itemMessage  = K.logStr msg
          , _itemTime     = itemTime
          , _itemNamespace = (env ^. KC.logEnvApp) <> (K.Namespace [ ns ])
          , _itemLoc      = Nothing

```

```

    }
    atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
else return False

```

Scribes

```

mkStdoutScribe :: K.Verbosity → IO K.Scribe
mkStdoutScribe = mkTextFileScribeH stdout True

mkStdoutScribeJson :: K.Verbosity → IO K.Scribe
mkStdoutScribeJson = mkJsonFileScribeH stdout True

mkStderrScribe :: K.Verbosity → IO K.Scribe
mkStderrScribe = mkTextFileScribeH stderr True

mkJsonFileScribeH :: Handle → Bool → K.Verbosity → IO K.Scribe
mkJsonFileScribeH handler color verb = do
    mkFileScribeH handler formatter color verb
where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter h _ verbosity item = do
        let tmsg = case KC._itemMessage item of
            K.LogStr "" → K.itemJson verbosity item
            K.LogStr msg → K.itemJson verbosity $
                item { KC._itemMessage = K.logStr (" " :: Text)
                    , KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
                    }
        TIO.hPutStrLn h (encodeToLazyText tmsg)

mkTextFileScribeH :: Handle → Bool → K.Verbosity → IO K.Scribe
mkTextFileScribeH handler color verb = do
    mkFileScribeH handler formatter color verb
where
    formatter h colorize verbosity item =
        TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item

mkFileScribeH
    :: Handle
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()) -- format and output f
    → Bool -- whether the output is colourized
    → K.Verbosity
    → IO K.Scribe
mkFileScribeH h formatter colorize verbosity = do
    hSetBuffering h LineBuffering
    locklocal ← newMVar ()
    let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
        logger item = withMVar locklocal $ \_ →
            formatter h colorize verbosity item
    pure $ K.Scribe logger (hClose h)

mkTextFileScribe :: Internal.FileDescription → Bool → Severity → K.Verbosity → IO K.Scribe

```

```

mkTextFileScribe fdesc colorize s v = do
  mkFileScribe fdesc formatter colorize s v
where
  formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
  formatter hdl colorize' v' item = do
    let tmsg = toLazyText $ formatItem colorize' v' item
    TIO.hPutStrLn hdl tmsg
mkFileScribe
  :: Internal.FileDescription
  → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()) -- format and output f
  → Bool -- whether the output is colourized
  → Severity
  → K.Verbosity
  → IO K.Scribe
mkFileScribe fdesc formatter colorize _ v = do
  let prefixDir = Internal.prefixPath fdesc
  (createDirectoryIfMissing True prefixDir)
  'catchIO' (Internal.prtoutException ("cannot log prefix directory: " ++ prefixDir))
  let fpath = Internal.filePath fdesc
  h ← catchIO (openFile fpath WriteMode) $
    λe → do
      Internal.prtoutException ("error while opening log: " ++ fpath) e
      -- fallback to standard output in case of exception
      return stdout
  hSetBuffering h LineBuffering
  scribestate ← newMVar h
  let finalizer :: IO ()
  finalizer = withMVar scribestate hClose
  let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
  logger item =
    withMVar scribestate $ λhandler →
      formatter handler colorize v item
  return $ K.Scribe logger finalizer

formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
  fromText header <>
  fromText " " <>
  brackets (fromText timestamp) <>
  fromText " " <>
  KC.unLogStr _itemMessage
where
  header = colorBySeverity _itemSeverity $
    "[ " <> mconcat namedcontext <> ": " <> severity <> ": " <> threadid <> " ]"
  namedcontext = KC.intercalateNs _itemNamespace
  severity = KC.renderSeverity _itemSeverity
  threadid = KC.getThreadIdText _itemThread

```



```

timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
tsformat :: String
tsformat = "%F %T%2Q %Z"
colorBySeverity s m = case s of
    K.EmergencyS → red m
    K.AlertS     → red m
    K.CriticalS  → red m
    K.ErrorS     → red m
    K.NoticeS    → magenta m
    K.WarningS   → yellow m
    K.InfoS      → blue m
    _            → m
red = colorize "31"
yellow = colorize "33"
magenta = colorize "35"
blue = colorize "34"
colorize c m
    | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
    | otherwise = m
-- translate Severity to Katip.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
    Debug → K.DebugS
    Info   → K.InfoS
    Notice → K.NoticeS
    Warning → K.WarningS
    Error  → K.ErrorS

```

1.4.13 Cardano.BM.Output.Aggregation

The aggregation is a singleton.

```

-- internal access to the aggregation
{-# NOINLINE aggregation #-}
aggregation :: MVar AggregationInternal
aggregation = unsafePerformIO $ do
    newMVar $ error "Aggregation MVar is not initialized."
-- Our internal state
data AggregationInternal = AggregationInternal
    { agMap :: HM.HashMap Text Aggregated
    , agSome :: [Int] -- TODO
    }

inspect :: Text → IO (Maybe Aggregated)
inspect name =
    withMVar aggregation $ λag →
        return $ HM.lookup name (agMap ag)

```

```

setup :: Configuration → IO ()
setup _ = do
  _ ← takeMVar aggregation
  -- TODO create thread which will periodically output
  -- aggregated values to the switchboard
  putMVar aggregation $ AggregationInternal HM.empty []

pass :: NamedLogItem → IO ()
pass item = do
  ag ← takeMVar aggregation
  putMVar aggregation $ AggregationInternal (updated $ agMap ag) (agSome ag)
where
  updated agmap = pass' (lnItem item) (lnName item) agmap
  pass' :: LogObject → LoggerName → HM.HashMap Text Aggregated → HM.HashMap Text Aggregated
  pass' (LP (LogValue iname value)) logname agmap =
    let name = logname <> "." <> iname
    in
      HM.alter ( $\lambda m \rightarrow \text{updateAggregation value } m$ ) name agmap
  -- TODO for text messages aggregate on delta of timestamps
  pass' _ _ agmap = agmap

```